

```
from google.colab import drive
drive.mount('/content/drive')
from warnings import simplefilter

simplefilter(action='ignore', category=FutureWarning)
```

➞ Go to this URL in a browser: https://accounts.google.com/o/oauth2/auth?client_id=947318989803-6bn6qk8qdgf4n4g3pfee649

Enter your authorization code:

.....

Mounted at /content/drive

▼ Homework 3 - Ames Housing Dataset

For all parts below, answer all parts as shown in the Google document for Homework 3. Be sure to include both code that justifies your answer as well as text to answer the questions. We also ask that code be commented to make it easier to follow.


▼ Part 1 - Pairwise Correlations

```
# TODO: show visualization
import pandas as pd
import numpy as np
trans = pd.read_csv('/content/drive/My Drive/Colab Notebooks/house-prices (1)/train.csv')
#trans.head()

trans = trans[['Id', 'LotFrontage', 'MSZoning', 'LotArea', 'Street', 'LandContour', 'LandSlope', 'Neighborhood', 'HouseStyle', 'OverallQual']]

trans.Neighborhood = pd.Categorical(trans.Neighborhood)
trans.Neighborhood = trans.Neighborhood.cat.codes
trans.LandSlope = pd.Categorical(trans.LandSlope)
trans.LandSlope = trans.LandSlope.cat.codes
trans.HouseStyle = pd.Categorical(trans.HouseStyle)
trans.HouseStyle = trans.HouseStyle.cat.codes
```

```
import seaborn as sns
sns.set(rc={'figure.figsize':(12,8)})
sns.heatmap(trans.corr(), vmin=-1, vmax=1, cmap='seismic', linewidths=0.2, annot=True);
```



Discuss most positive and negative correlations.


The most positive correlations with SalePrice are OverallQual, GrLivArea, GarageCars, YearBuilt. OverallQual correlates positively with GarageCars and GrLivArea. There are not many negatively correlated fields, housestyle correlates negatively with LotArea. Landslope correlates negatively with YearBuilt, OverallQual, YearRemodAdd.

▼ Part 2 - Informative Plots

```
# TODO: code to generate Plot 1
import seaborn as sns
trans = pd.read_csv('/content/drive/My Drive/Colab Notebooks/house-prices (1)/train.csv')

sns.set(rc={'figure.figsize':(8,4)})
fig=sns.barplot(trans['HouseStyle'].unique()[0:10],trans['HouseStyle'].value_counts()[0:10])
fig.set(xlabel='HouseStyle', ylabel='Count of houses')

[Text(0, 0.5, 'Count of houses'), Text(0.5, 0, 'HouseStyle')]
```




What interesting properties does Plot 1 reveal?

Most houses were 2 storey, followed by 1 storey.

```
sns.set(rc={'figure.figsize':(12,6)})
fig=sns.lineplot(trans['YearBuilt'].unique()[0:1000],trans['YearBuilt'].value_counts()[0:1000])
```

```
fig.set(xlabel='YearBuilt', ylabel='Number of houses built')
```

```
↳ [Text(0, 0.5, 'Number of houses built'), Text(0.5, 0, 'YearBuilt')]
```




What interesting properties does Plot 2 reveal?

The highest number of houses were built in the year 2003 and 1976.

```
import numpy as np
fig=sns.scatterplot((trans['GrLivArea']),trans['SalePrice'])
fig.set(xlabel='GrLivArea', ylabel='SalePrice')
```

↳

```
[Text(0, 0.5, 'SalePrice'), Text(0.5, 0, 'GrLivArea')]
```



What interesting properties does Plot 3 reveal?

The SalePrice appears to increase with increase in GrLiving area for most cases. Most houses have GrLiving area between 1000 and 2000 units.


```
sns.set(rc={'figure.figsize':(8,4)})

group = trans.groupby(['OverallQual'])
group = group.mean()
group = group.reset_index();

fig=sns.barplot(group['OverallQual'].unique(),group['SalePrice'])
fig.set(xlabel='OverallQual', ylabel='Average SalePrice')
```



```
[Text(0, 0.5, 'Average SalePrice'), Text(0.5, 0, 'OverallQual')]
```



What interesting properties does Plot 4 reveal?

Average Saleprice increases exponentially with Overall quality.

```
# trans['Neighborhood'].value_counts()
sns.set(rc={'figure.figsize':(22,6)})

group = trans.groupby(['Neighborhood'])
group = group.mean()
group = group.reset_index();
group = group.sort_values(by=['SalePrice'],ascending=False)
# temp5 = group['Neighborhood']


fig.set(xlabel='Neighborhood', ylabel='Avg SalePrice')
print("Most Expensive Neighborhoods: NoRidge, NRidgHt,StonBr")
print("Least Expensive Neighborhoods: MeadowV, IDOTRR, BrDale")

fig=sns.lineplot(group['Neighborhood'][:12], group['SalePrice'][:12]) #Most expensive
print('\n')
fig1=sns.lineplot(group['Neighborhood'][13:], group['SalePrice'][13:]) #Least expensive
```



Most Expensive Neighborhoods: NoRidge, NRidgHt, StonBr

Least Expensive Neighborhoods: MeadowV, IDOTRR, BrDale



What interesting properties does Plot 5 reveal?

Most Expensive Neighborhoods: NoRidge, NRidgHt, StonBr

Least Expensive Neighborhoods: MeadowV, IDOTRR, BrDale

▼ Part 3 - Handcrafted Scoring Function

TODO: code for scoring function

```

dep_vars = trans[['Id','Neighborhood', 'OverallQual','YearBuilt','SalePrice']]
dep_vars = dep_vars.reset_index();

dep_vars['score'] = ((dep_vars['OverallQual']*1000 + dep_vars['YearBuilt']*10 + dep_vars['SalePrice']/10)/10000)    #score
maxscore = max(dep_vars['score'])    # normalize scores
dep_vars['score'] = dep_vars['score']/maxscore    # score out of 1

print('MOST DESIRABLE')
print(dep_vars.sort_values(by=['score'],ascending=False).head(10))
print('\n')
print('LEAST DESIRABLE')
print(dep_vars.sort_values(by=['score']).head(10))

```

☞ MOST DESIRABLE

	index	Id	Neighborhood	OverallQual	YearBuilt	SalePrice	score
691	691	692	NoRidge	10	1994	755000	1.000000
1182	1182	1183	NoRidge	10	1996	745000	0.990706
1169	1169	1170	NoRidge	10	1995	625000	0.876802
898	898	899	NridgHt	9	2009	611657	0.855991
803	803	804	NridgHt	9	2008	582933	0.828654
440	440	441	NridgHt	10	2008	555000	0.811646
1046	1046	1047	StoneBr	9	2005	556581	0.803377
769	769	770	StoneBr	8	2003	538000	0.776081
178	178	179	StoneBr	9	2008	501837	0.751742
798	798	799	NridgHt	9	2008	485000	0.735774

LEAST DESIRABLE

	index	Id	Neighborhood	OverallQual	YearBuilt	SalePrice	score
533	533	534	BrkSide	1	1946	39300	0.231316
916	916	917	IDOTRR	2	1949	35311	0.237302
968	968	969	OldTown	3	1910	37900	0.245542
375	375	376	Edwards	1	1922	61000	0.249621
495	495	496	IDOTRR	4	1920	34900	0.253130
1100	1100	1101	SWISU	2	1920	60000	0.257967
30	30	31	IDOTRR	4	1920	40000	0.257967
636	636	637	BrkSide	2	1936	60000	0.259484
710	710	711	BrkSide	3	1935	52000	0.261286
1380	1380	1381	Edwards	3	1914	58500	0.265459

What is the ten most desirable houses?

(Listed Above)

What is the ten least desirable houses?

(Listed Above)

Describe your scoring function and how well you think it worked.

The score is dependent on 3 variables - Overall quality, Year built and SalePrice. All of these are positively correlated to desirability. Higher the quality, newer the house construction and higher the price - more desirable the house is.

▼ Part 4 - Pairwise Distance Function

```
# TODO: code for distance function
from sklearn.metrics.pairwise import euclidean_distances
trans.head()
pairwise = trans[['OverallQual', 'YearBuilt', 'SalePrice', 'GarageCars']]
# pairwise.head()
distances = euclidean_distances(pairwise, pairwise)
#pd.DataFrame(distances).head(10).hist()
max_ele = max(map(max, distances))
min_ele = np.amin(distances[distances != np.amin(distances)])
# print(min_ele)

print('Max distance is ', max(map(max, distances)))

list=[]
for i in range(0,1459):
    for j in range(1,1459):
        if distances[i][j] == max_ele:
            if i>=j:
                print("Max distance between house numbers ", (" ", i+1, " ", j+1, ' '))

print('Min distance is ', np.amin(distances[distances != np.amin(distances)]))

for i in range(0, 1459):
    for j in range(0, 1459):
        if distances[i][j] == 1.0:
            if i>=j:

                list.append((' '+str(i+1)+' ', '+str(j+1)+' '))
```

```
print('Min distance between house numbers ', list)
```

```

↳ Max distance is 720100.0038334953
   Max distance between house numbers ( 692 , 496 )
   Min distance is 1.0
   Min distance between house numbers ['(194, 146)', '(258, 163)', '(274, 20)', '(332, 274)', '(365, 110)', '(367, 287)']

```

```

print('\n', 'Eg. for max dist:')
print(trans[(trans.Id == 692) | (trans.Id == 496)])
print('\n', 'Eg. for min dist:')
print(trans[(trans.Id == 194) | (trans.Id == 146)])

```

```

↳
   Eg. for max dist:
      Id  MSSubClass  MSZoning  ...  SaleType  SaleCondition  SalePrice
495  496           30    C (all)  ...      WD      Abnorml      34900
691  692           60      RL    ...      WD      Normal      755000

```

[2 rows x 81 columns]

```

   Eg. for min dist:
      Id  MSSubClass  MSZoning  ...  SaleType  SaleCondition  SalePrice
145  146           160      RM    ...      WD      Normal      130000
193  194           160      RM    ...      WD      Normal      130000

```

[2 rows x 81 columns]

How well does the distance function work? When does it do well/badly?

The distance vector is max for properties 496 and 692. Looking at fields 'OverallQual','YearBuilt','SalePrice','GarageCars' we observe that they are extremely different in values.

The distance vector is min for a number of properties. For eg 194 and 146. Looking at fields 'OverallQual','YearBuilt','SalePrice','GarageCars' we observe that they are quite similar/same in values.

▼ Part 5 - Clustering

```

# TODO: code for clustering and visualization
# aggllo
from sklearn.cluster import KMeans
import numpy as np

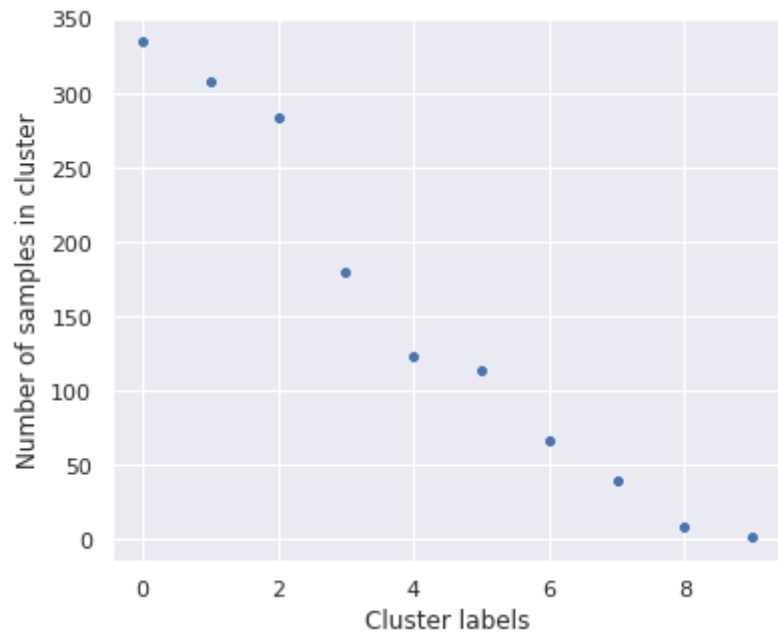
kmeans = KMeans(n_clusters=10, random_state=0).fit(distances)
kmeans.labels_
kmeans.predict(distances)
arr = kmeans.labels_
df=trans[['Id','Neighborhood']].copy()
df['labels'] = arr

fig=sns.scatterplot(np.unique(arr),pd.Series(arr).value_counts())
sns.set(rc={'figure.figsize':(6,5)})

fig.set(xlabel='Cluster labels', ylabel='Number of samples in cluster')
print('Cluster labels:')
arr

```


Cluster labels:
array([6, 3, 6, ..., 1, 9, 9], dtype=int32)



```
sns.set(rc={'figure.figsize':(22,6)})

fig=sns.lineplot((df['Neighborhood']),df['labels'])
fig.set(xlabel='Neighborhoods', ylabel='Cluster labels (0-9)')

plt.annotate([Text(0, 0.5, 'Cluster labels (0-9)'), Text(0.5, 0, 'Neighborhoods')])
```



How well do the clusters reflect neighborhood boundaries? Write a discussion on what your clusters capture and how well they work.

The distance matrix measures euclidean distances of variables 'OverallQual','YearBuilt','SalePrice','GarageCars'. The clusters reflect neighborhood boundaries to a good extent for almost all neighborhoods(with exceptions of Blueste, NPKvill, Veenkar as shown in the graph).

▼ Part 6 - Linear Regression

```
# TODO: code for linear regression
import pandas as pd
```

```
from sklearn.linear_model import LogisticRegression, LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import CountVectorizer
from sklearn import metrics
from sklearn.metrics import mean_squared_error
from sklearn.metrics import mean_squared_log_error
from math import sqrt
from sklearn import preprocessing
import matplotlib.pyplot as plt

trans = pd.read_csv('/content/drive/My Drive/Colab Notebooks/house-prices (1)/train.csv')

trans_X = trans[['OverallQual', 'YearBuilt', 'SalePrice', 'GarageCars', 'GrLivArea']]
trans_X = trans_X.fillna(trans_X.mode().iloc[0])

trans_y = pd.DataFrame(trans_X['SalePrice'])

trans_X = trans_X.drop(['SalePrice'], axis=1)

X, y = trans_X, trans_y
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=500)

clf = LinearRegression()
clf.fit(X_train, y_train)

preds = clf.predict(X_test)
print("MEAN SQUARED LOG ERROR: ", np.sqrt(metrics.mean_squared_log_error(preds, y_test)))
#-----


test = pd.read_csv('/content/drive/My Drive/Colab Notebooks/house-prices (1)/test.csv')
test = test[['OverallQual', 'YearBuilt', 'GarageCars', 'GrLivArea']]
test = test.fillna(test.mode().iloc[0])

preds = clf.predict(test)
plt.hist(preds)
plt.xlabel('Sale Price predictions')
plt.ylabel('count')

plt.show()
```



MEAN SQUARED LOG ERROR: 0.24335448037635204



How well/badly does it work? Which are the most important variables?

It worked well - RMSLE = 0.24. The most important variables were 'OverallQual', 'YearBuilt', 'GarageCars', 'GrLivArea'. These are the ones with a high correlation with SalePrice.

▼ Part 7 - External Dataset

TODO: code to import external dataset and test

```
inflation = pd.read_csv('/content/inflation.csv') #Average CPI for years 1913-2018
inflation['YrSold'] = inflation['Year']
inflation=inflation.drop(['Year'],axis=1)
```

```
trans = pd.read_csv('/content/drive/My Drive/Colab Notebooks/house-prices (1)/train.csv')
# trans_X=trans[['Id','LotFrontage','LotArea','OverallQual','OverallCond','YearBuilt','YearRemodAdd','GarageCars','PoolArea','YrSold']]

trans_X = trans[['OverallQual','YearBuilt','SalePrice','GarageCars','GrLivArea','YrSold']]
trans_X = pd.merge(trans_X, inflation, on='YrSold', how='left', left_index=True)
```

```

trans_X

trans_X = trans_X.fillna(trans_X.mode().iloc[0])

trans_y = pd.DataFrame(trans_X['SalePrice'])

trans_X = trans_X.drop(['SalePrice'], axis=1)

X, y = trans_X, trans_y
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=500)

clf = LinearRegression()
clf.fit(X_train, y_train)

preds = clf.predict(X_test)
print("MEAN SQUARED LOG ERROR: ", np.sqrt(metrics.mean_squared_log_error(preds, y_test)))

➤ MEAN SQUARED LOG ERROR: 0.24638999666977993

```

Describe the dataset and whether this data helps with prediction.

The dataset used is called 'Historical Consumer Price Index (CPI-U) Data' from <https://inflationdata.com>. The Consumer Price Index (CPI) is a measure of the average change over time in the prices paid by urban consumers for a market basket of consumer goods and services. A Consumer Price Index of 158 indicates 58% inflation since 1982, while a CPI index of 239 would indicate 139% inflation since 1982. CPI has a high correlation with property rates, hence it should improve our model accuracy.

▼ Part 8 - Permutation Test

```

# TODO: code for all permutation tests
from mlxtend.evaluate import permutation_test
from sklearn.model_selection import permutation_test_score

import pandas as pd
from sklearn.linear_model import LogisticRegression, LinearRegression
from sklearn.ensemble import AdaBoostRegressor, RandomForestRegressor
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import CountVectorizer
from sklearn import metrics
from sklearn.metrics import mean_squared_error
from sklearn.metrics import mean_squared_log_error

from math import sqrt
from sklearn import preprocessing

```

```

trans = pd.read_csv('/content/drive/My Drive/Colab Notebooks/house-prices (1)/train.csv')
# trans_X=trans[['Id','LotFrontage','LotArea','OverallQual','OverallCond','YearBuilt','YearRemodAdd','GarageCars','PoolArea','YrSold']]

trans_X=trans[['OverallQual','SalePrice']]
trans_X = trans_X.fillna(trans_X.mode().iloc[0])

trans_y = (trans_X['SalePrice'])

trans_X = trans_X.drop(['SalePrice'], axis=1)

X, y = trans_X, trans_y
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=500)

clf = RandomForestRegressor(n_estimators=100, random_state=20)
model = clf.fit(X_train, y_train)
preds = clf.predict(X_test)

score, permutation_scores, pvalue = permutation_test_score(model, trans_X, trans_y, n_permutations=100, n_jobs=1)
print('Pvalue for OverallQual', pvalue)
print('RMSLE', np.sqrt(metrics.mean_squared_log_error(preds, y_test)))
print('\n')
#-----
trans_X=trans[['YearBuilt','SalePrice']]
trans_X = trans_X.fillna(trans_X.mode().iloc[0])

trans_y = (trans_X['SalePrice'])

trans_X = trans_X.drop(['SalePrice'], axis=1)

X, y = trans_X, trans_y
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=500)

clf = RandomForestRegressor(n_estimators=100, random_state=20)
model = clf.fit(X_train, y_train)
preds = clf.predict(X_test)

score, permutation_scores, pvalue = permutation_test_score(model, trans_X, trans_y, n_permutations=100, n_jobs=1)
print('Pvalue for YearBuilt', pvalue)
print('RMSLE', np.sqrt(metrics.mean_squared_log_error(preds, y_test)))
print('\n')
# #-----

trans_X=trans[['GarageCars','SalePrice']]
trans_X = trans_X.fillna(trans_X.mode().iloc[0])

trans_y = (trans_X['SalePrice'])

trans_X = trans_X.drop(['SalePrice'], axis=1)

```



```
X, y = trans_X, trans_y
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=500)

clf = RandomForestRegressor(n_estimators=100, random_state=20)
model = clf.fit(X_train, y_train)
preds = clf.predict(X_test)

score, permutation_scores, pvalue = permutation_test_score(model, trans_X, trans_y, n_permutations=100, n_jobs=1)
print('Pvalue for GarageCars', pvalue)
print('RMSLE', np.sqrt(metrics.mean_squared_log_error(preds, y_test)))
print('\n')
```

```
# #-----

trans_X=trans[['GrLivArea', 'SalePrice']]
trans_X = trans_X.fillna(trans_X.mode().iloc[0])

trans_y = (trans_X['SalePrice'])

trans_X = trans_X.drop(['SalePrice'], axis=1)

X, y = trans_X, trans_y
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=500)

clf = RandomForestRegressor(n_estimators=100, random_state=20)
model = clf.fit(X_train, y_train)
preds = clf.predict(X_test)

score, permutation_scores, pvalue = permutation_test_score(model, trans_X, trans_y, n_permutations=100, n_jobs=1)
print('Pvalue for GrLivArea', pvalue)
print('RMSLE', np.sqrt(metrics.mean_squared_log_error(preds, y_test)))
print('\n')
```

```
# #-----

trans_X=trans[['Id', 'SalePrice']]
trans_X = trans_X.fillna(trans_X.mode().iloc[0])

trans_y = (trans_X['SalePrice'])

trans_X = trans_X.drop(['SalePrice'], axis=1)

X, y = trans_X, trans_y
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=500)

clf = RandomForestRegressor(n_estimators=100, random_state=20)
model = clf.fit(X_train, y_train)
preds = clf.predict(X_test)
```

```
score, permutation_scores, pvalue = permutation_test_score(model, trans_X, trans_y, n_permutations=100, n_jobs=1)
print('Pvalue for Id', pvalue)
print('RMSLE', np.sqrt(metrics.mean_squared_log_error(preds, y_test)))
print('\n')
```

```
# #-----

trans_X=trans[['BsmtFinSF2', 'SalePrice']]
trans_X = trans_X.fillna(trans_X.mode().iloc[0])

trans_y = (trans_X['SalePrice'])

trans_X = trans_X.drop(['SalePrice'], axis=1)

X, y = trans_X, trans_y
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=500)

clf = RandomForestRegressor(n_estimators=100, random_state=20)
model = clf.fit(X_train, y_train)

score, permutation_scores, pvalue = permutation_test_score(model, trans_X, trans_y, n_permutations=100, n_jobs=1)
print('Pvalue for BsmtFinSF2', pvalue)
print('RMSLE', np.sqrt(metrics.mean_squared_log_error(preds, y_test)))
print('\n')
```

```
# #-----

trans_X=trans[['PoolArea', 'SalePrice']]
trans_X = trans_X.fillna(trans_X.mode().iloc[0])

trans_y = (trans_X['SalePrice'])

trans_X = trans_X.drop(['SalePrice'], axis=1)

X, y = trans_X, trans_y
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=500)

clf = RandomForestRegressor(n_estimators=100, random_state=20)
model = clf.fit(X_train, y_train)
preds = clf.predict(X_test)

score, permutation_scores, pvalue = permutation_test_score(model, trans_X, trans_y, n_permutations=100, n_jobs=1)
print('Pvalue for PoolArea', pvalue)
print('RMSLE', np.sqrt(metrics.mean_squared_log_error(preds, y_test)))
print('\n')
```

```
# #-----

trans_X=trans[['OverallCond', 'SalePrice']]
trans_X = trans_X.fillna(trans_X.mode().iloc[0])
```

```
trans_y = (trans_X['SalePrice'])

trans_X = trans_X.drop(['SalePrice'], axis=1)

X, y = trans_X, trans_y
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=500)

clf = RandomForestRegressor(n_estimators=100, random_state=20)
model = clf.fit(X_train, y_train)
preds = clf.predict(X_test)

score, permutation_scores, pvalue = permutation_test_score(model, trans_X, trans_y, n_permutations=100, n_jobs=1)
print('Pvalue for OverallCond', pvalue)
print('RMSLE', np.sqrt(metrics.mean_squared_log_error(preds, y_test)))
print('\n')
```

#-----

```
trans_X=trans[['TotalBsmtSF', 'SalePrice']]
trans_X = trans_X.fillna(trans_X.mode().iloc[0])

trans_y = (trans_X['SalePrice'])

trans_X = trans_X.drop(['SalePrice'], axis=1)

X, y = trans_X, trans_y
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=500)

clf = RandomForestRegressor(n_estimators=100, random_state=20)
model = clf.fit(X_train, y_train)
preds = clf.predict(X_test)

score, permutation_scores, pvalue = permutation_test_score(model, trans_X, trans_y, n_permutations=100, n_jobs=1)
print('Pvalue for TotalBsmtSF', pvalue)
print('RMSLE', np.sqrt(metrics.mean_squared_log_error(preds, y_test)))
print('\n')
```

#-----

```
trans_X=trans[['KitchenAbvGr', 'SalePrice']]
trans_X = trans_X.fillna(trans_X.mode().iloc[0])

trans_y = (trans_X['SalePrice'])

trans_X = trans_X.drop(['SalePrice'], axis=1)

X, y = trans_X, trans_y
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=500)

clf = RandomForestRegressor(n_estimators=100, random_state=20)
```

```
model = clf.fit(X_train, y_train)
preds = clf.predict(X_test)

score, permutation_scores, pvalue = permutation_test_score(model, trans_X, trans_y, n_permutations=100, n_jobs=1)
print('Pvalue for KitchenAbvGr', pvalue)
print('RMSLE', np.sqrt(metrics.mean_squared_log_error(preds, y_test)))
print('\n')
```

#-----



Pvalue for OverallQual 0.009900990099009901
RMSLE 0.22348276760110886

Pvalue for YearBuilt 0.009900990099009901
RMSLE 0.2863191565877305

Pvalue for GarageCars 0.009900990099009901
RMSLE 0.26854153922440266

Pvalue for GrLivArea 0.009900990099009901
RMSLE 0.2974631532054856

Pvalue for Id 0.09900990099009901
RMSLE 0.44984043989531625

Pvalue for BsmtFinSF2 0.019801980198019802
RMSLE 0.44984043989531625

Pvalue for PoolArea 0.019801980198019802
RMSLE 0.3675701104872097

Pvalue for OverallCond 0.009900990099009901
RMSLE 0.3347009508097016

Pvalue for TotalBsmtSF 0.009900990099009901
RMSLE 0.3274647248042461

Pvalue for KitchenAbvGr 0.009900990099009901
RMSLE 0.3656845581892145

Describe the results.

The p-value is used in the context of null hypothesis testing in order to quantify the idea of statistical significance of evidence. A small p-value (typically ≤ 0.05) indicates strong evidence against the null hypothesis, so you reject the null hypothesis. The p-values here are less than 5%(statistically significant) indicating that the Random Forest Regression model is a good choice.

▼ Part 9 - Final Result

```
# TODO: code for RandomForestRegressor
import pandas as pd
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import CountVectorizer
from sklearn import metrics
from sklearn.metrics import mean_squared_error
from math import sqrt
from sklearn import preprocessing
import xgboost as xgb

trans = pd.read_csv('/content/drive/My Drive/Colab Notebooks/house-prices (1)/train.csv')
# trans_X=trans[['Id','LotFrontage','LotArea','OverallQual','OverallCond','YearBuilt','YearRemodAdd','GarageCars','PoolArea','YrSold']]

trans_X=trans[['OverallQual','YearBuilt','SalePrice','GarageCars','GrLivArea']]
trans_X = trans_X.fillna(trans_X.mode().iloc[0])

trans_y = pd.DataFrame(trans_X['SalePrice'])

trans_X = trans_X.drop(['SalePrice'], axis=1)

X, y = trans_X, trans_y
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=500)

clf = xgb.XGBRegressor(learning_rate=0.01,n_estimators=3460,
                        max_depth=3, min_child_weight=0,
                        gamma=0, subsample=0.7,
                        colsample_bytree=0.7,
                        objective='reg:linear', nthread=-1,
                        scale_pos_weight=1, seed=27,
                        reg_alpha=0.00006)
```

XGBoost stands for "Extreme Gradient Boosting" is a supervised machine learning model, rmsle obtained = 0.1542

```
var = clf.fit(X_train, y_train)
```

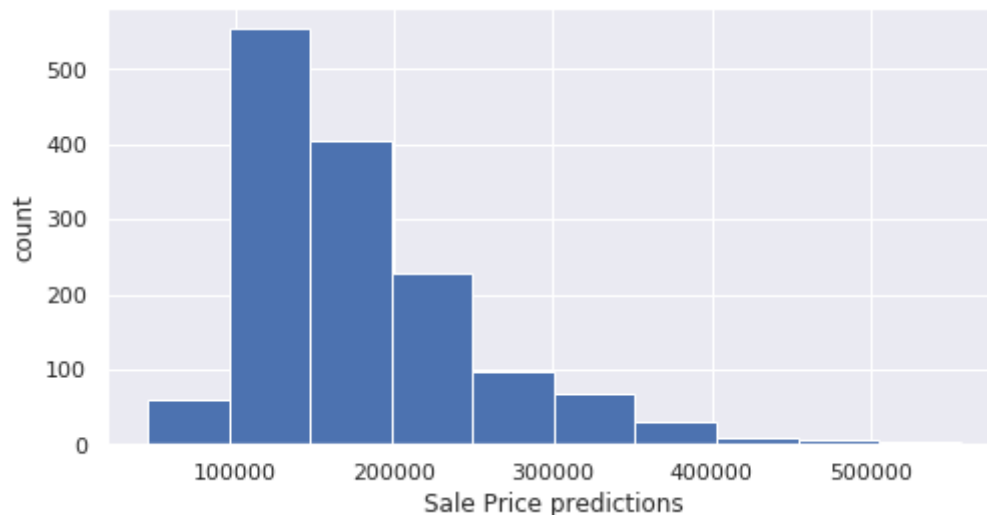
```
preds = clf.predict(X_test)
print('RMSLE: ', np.sqrt(metrics.mean_squared_log_error(preds, y_test)))
```

```
test = pd.read_csv('/content/drive/My Drive/Colab Notebooks/house-prices (1)/test.csv')
test = test[['OverallQual', 'YearBuilt', 'GarageCars', 'GrLivArea']]
test = test.fillna(test.mode().iloc[0])
```

```
preds = clf.predict(test)
# pd.DataFrame(preds).hist()
plt.hist(preds)
plt.xlabel('Sale Price predictions')
plt.ylabel('count')
```

```
plt.show()
```

⏏ [10:46:24] WARNING: /workspace/src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of reg:squarederror
RMSLE: 0.1542517704280983



```
pd.DataFrame(preds).to_csv("output03.csv")
```

▼ Part 10 - Kaggle Submission

Report the rank, score, number of entries, for your highest rank. Include a snapshot of your best score on the leaderboard as confirmation. Be sure to provide a link to your Kaggle profile. Make sure to include a screenshot of your ranking. Make sure your profile includes your face and affiliation with SBU.

Kaggle Link: <https://www.kaggle.com/saylik>

Highest Rank: 3781

Score: 0.17506

Number of entries: 5

INCLUDE IMAGE OF YOUR KAGGLE RANKING: <https://drive.google.com/file/d/13ZGrRwrGMExtcKK10wgBnGvYO6E9mS28/view?usp=sharing>

